

PageRank optimization

```
lines = sc.textFile("links.txt")
links = lines.map(parseNeighbors)
            .distinct()
            .groupByKey()
            .cache()
ranks = links.map(lambda url_neighbors: (url_neighbors[0], 1.0))


for iteration in xrange(n):
    contribs = links.join(ranks).flatMap(
        lambda u: computeContribs(u[1][0], u[1][1]))
    ranks = contribs.reduceByKey(add)
                .mapValues(lambda rank: rank * 0.85 + 0.15)

ranks.collect()
```

```
lines = sc.textFile("links.txt")
links = lines.map(parseNeighbors)
               .distinct()
               .groupByKey()
               .cache()
ranks = links.map(lambda url_neighbors: (url_neighbors[0], 1.0))

for iteration in xrange(n):
    contribs = links.join(ranks).flatMap(
        lambda u: computeContribs(u[1][0], u[1][1]))
    ranks = contribs.reduceByKey(add)
                  .mapValues(lambda rank: rank * 0.85 + 0.15)

ranks.collect()
```



- 10 executors, 2 cores & 4GB per executor
- 9 lines
- 30 iterations

Spark Jobs (?)

User: pklemenkov

Total Uptime: 1.1 min

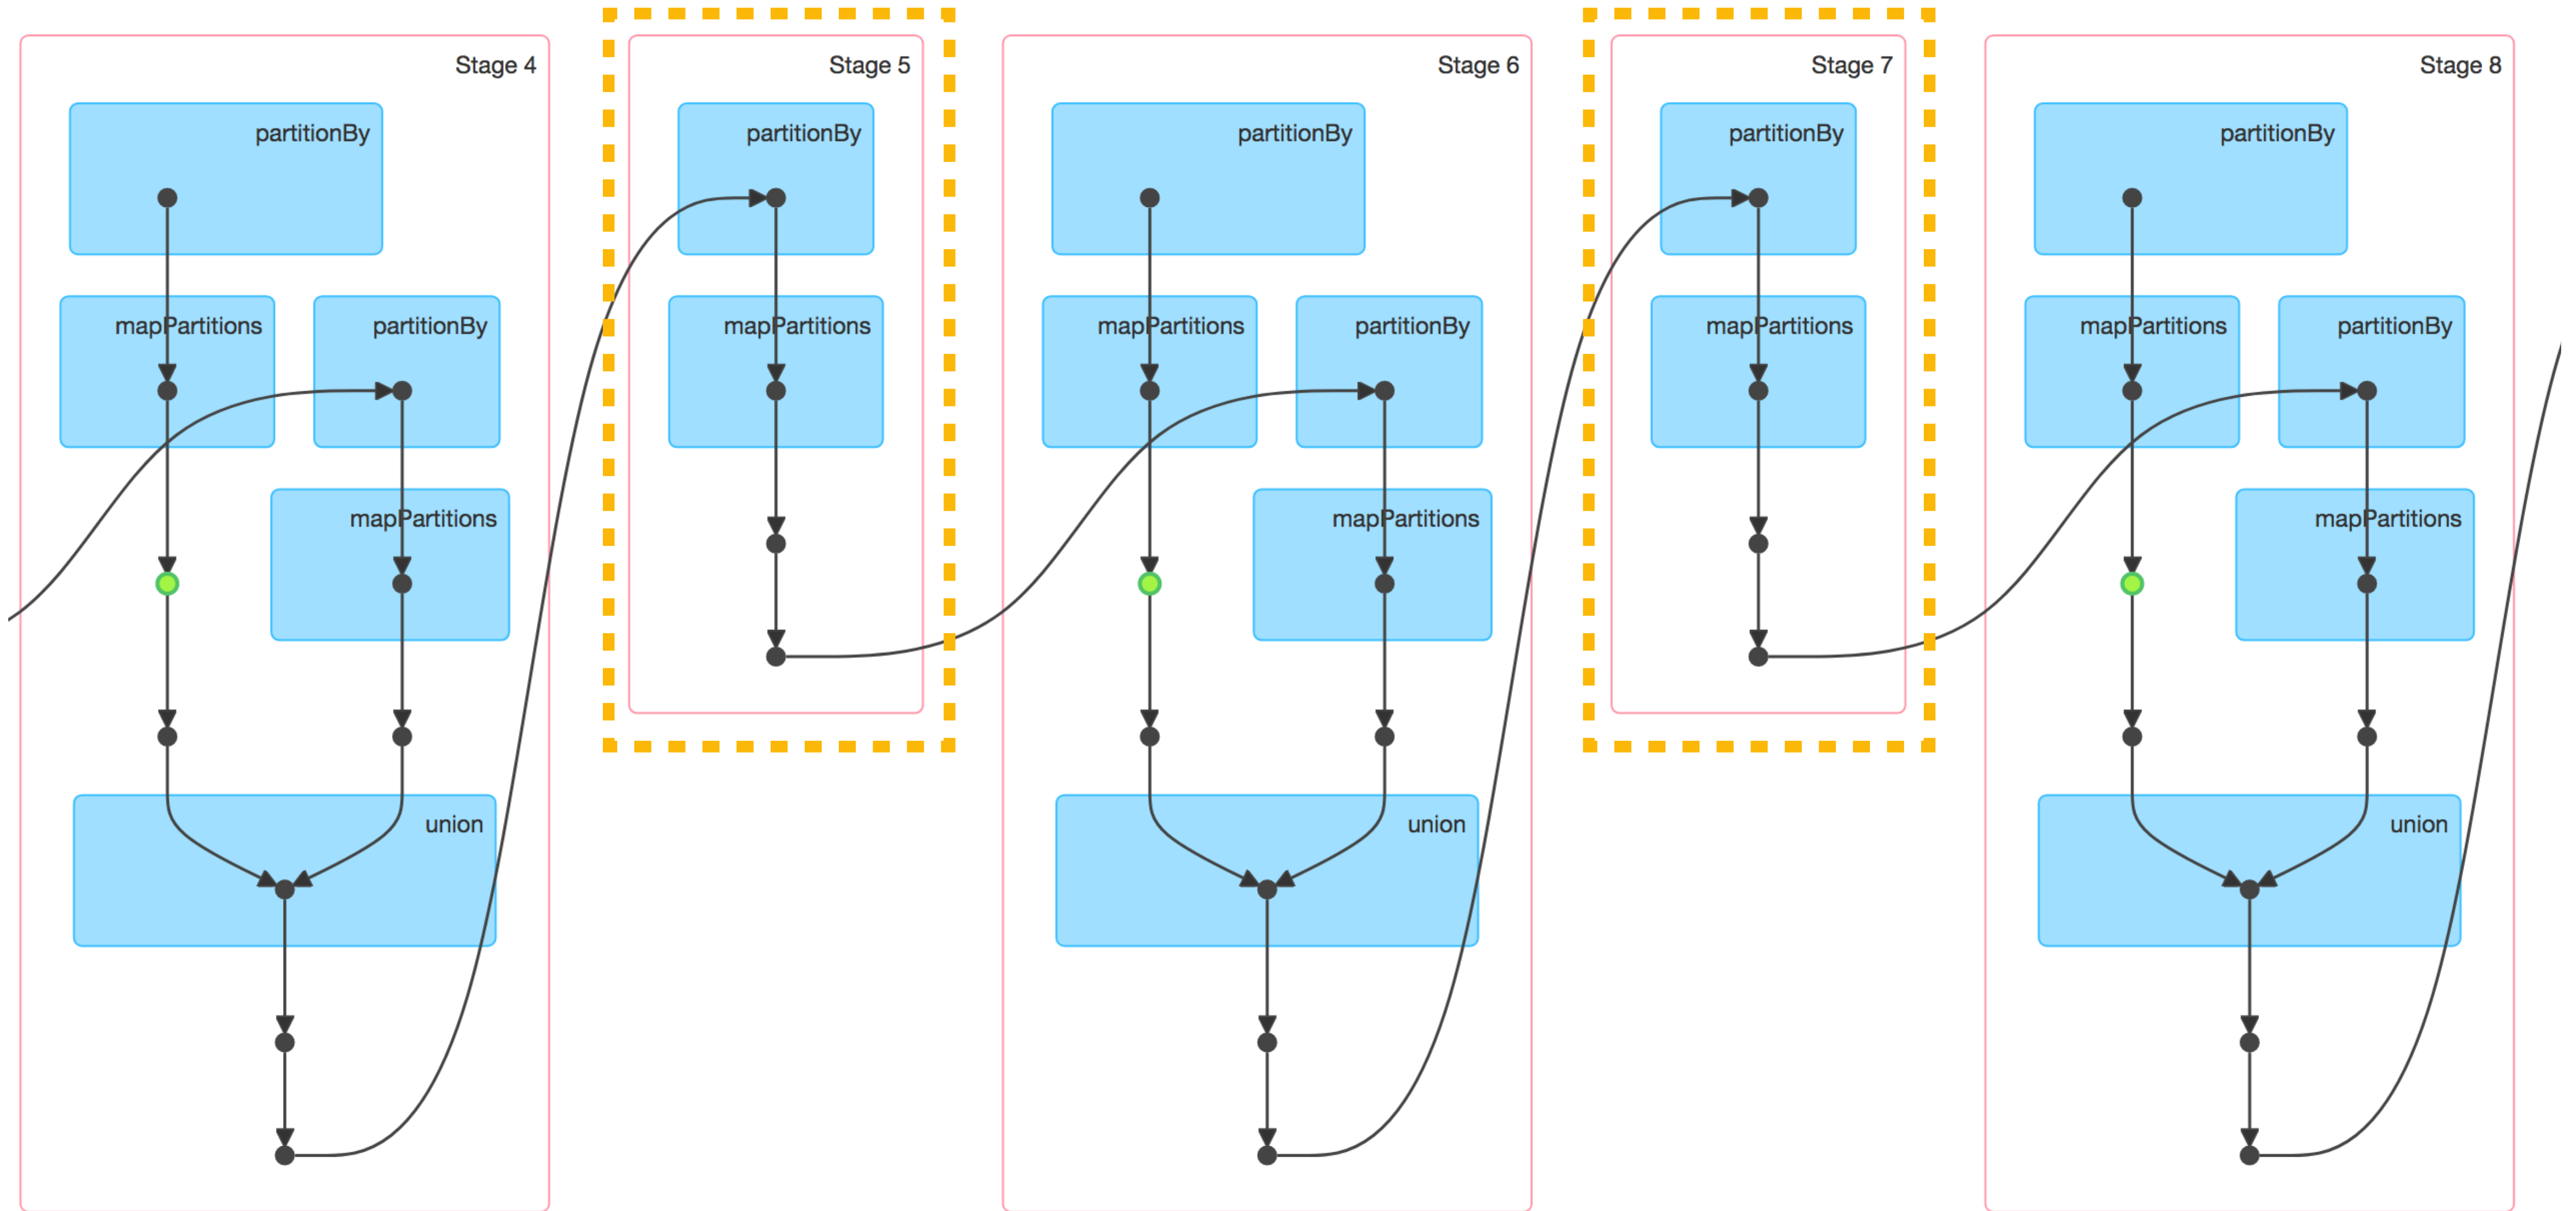
Scheduling Mode: FIFO

Completed Jobs: 1

▶ Event Timeline

Completed Jobs (1)

Job Id ▾	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
0	collect at <magic-timeit>:257	2017/07/25 04:50:52	24 s	63/63	2046/2046



```
for iteration in xrange(n):  
    contribs = links.join(ranks).flatMap(  
        lambda u: computeContribs(u[1][0], u[1][1]))  
    ranks = contribs.reduceByKey(add)  
        .mapValues(lambda rank: rank * 0.85 + 0.15)
```



```
links = lines.map(parseNeighbors)
            .distinct()
            .groupByKey()
            .cache()
ranks = links.map(lambda url_neighbors: (url_neighbors[0], 1.0))

>> links.getNumPartitions(), links.partitioner
(2, <pyspark.rdd.Partitioner at 0x7f1ffdb65050>)

>> ranks.getNumPartitions(), ranks.partitioner
(2, None)

>> links.partitioner == ranks.partitioner
False
```

```
links = lines.map(parseNeighbors)
            .distinct()
            .groupByKey()
            .partitionBy(2)
            .cache()

ranks = links.map(lambda url_neighbors: (url_neighbors[0], 1.0))

>> links.getNumPartitions(), links.partitioner
(2, <pyspark.rdd.Partitioner at 0x7f1ffdb65050>)

>> ranks.getNumPartitions(), ranks.partitioner
(2, None)

>> links.partitioner == ranks.partitioner
False
```

```
links = lines.map(parseNeighbors)
            .distinct()
            .groupByKey()
            .partitionBy(2)
            .cache()
ranks = links.map(lambda url_neighbors: (url_neighbors[0], 1.0))

>> links.getNumPartitions(), links.partitioner
(2, <pyspark.rdd.Partitioner at 0x7f1ffdb65050>)

>> ranks.getNumPartitions(), ranks.partitioner
(2, None)

>> links.partitioner == ranks.partitioner
False
```

```
links = lines.map(parseNeighbors)
            .distinct()
            .groupByKey()
            .partitionBy(2)
            .cache()
ranks = links.map(lambda url_neighbors: (url_neighbors[0], 1.0),
                preservePartitioning=True)
```

```
>> links.getNumPartitions(), links.partitioner
(2, <pyspark.rdd.Partitioner at 0x7f1ffdb65050>)
```

```
>> ranks.getNumPartitions(), ranks.partitioner
(2, <pyspark.rdd.Partitioner at 0x7f1ffdb65050>)
```

```
>> links.partitioner == ranks.partitioner
True
```

```
for iteration in xrange(n):  
    contribs = links.join(ranks).flatMap(  
        lambda u: computeContribs(u[1][0], u[1][1]))  
    ranks = contribs.reduceByKey(add,  
                                numPartitions=links.getNumPartitions())  
    .mapValues(lambda rank: rank * 0.85 + 0.15)
```

```
for iteration in xrange(n):  
    contribs = links.join(ranks).flatMap(  
        lambda u: computeContribs(u[1][0], u[1][1]))  
    ranks = contribs.reduceByKey(add,  
                                numPartitions=links.getNumPartitions())  
    .mapValues(lambda rank: rank * 0.85 + 0.15)
```

```
links = lines.map(parseNeighbors)
            .distinct()
            .groupByKey()
            .partitionBy(2)
            .cache()
ranks = links.map(lambda url_neighbors: (url_neighbors[0], 1.0)),
            preservePartitioning=True)

for iteration in xrange(n):
    contribs = links.join(ranks).flatMap(
        lambda u: computeContribs(u[1][0], u[1][1]))
    ranks = contribs.reduceByKey(add,
                                numPartitions=links.getNumPartitions())
        .mapValues(lambda rank: rank * 0.85 + 0.15)
ranks.collect()
```

Spark Jobs (?)


User: pklemenkov
Total Uptime: 55 s
Scheduling Mode: FIFO
Completed Jobs: 1

▶ [Event Timeline](#)

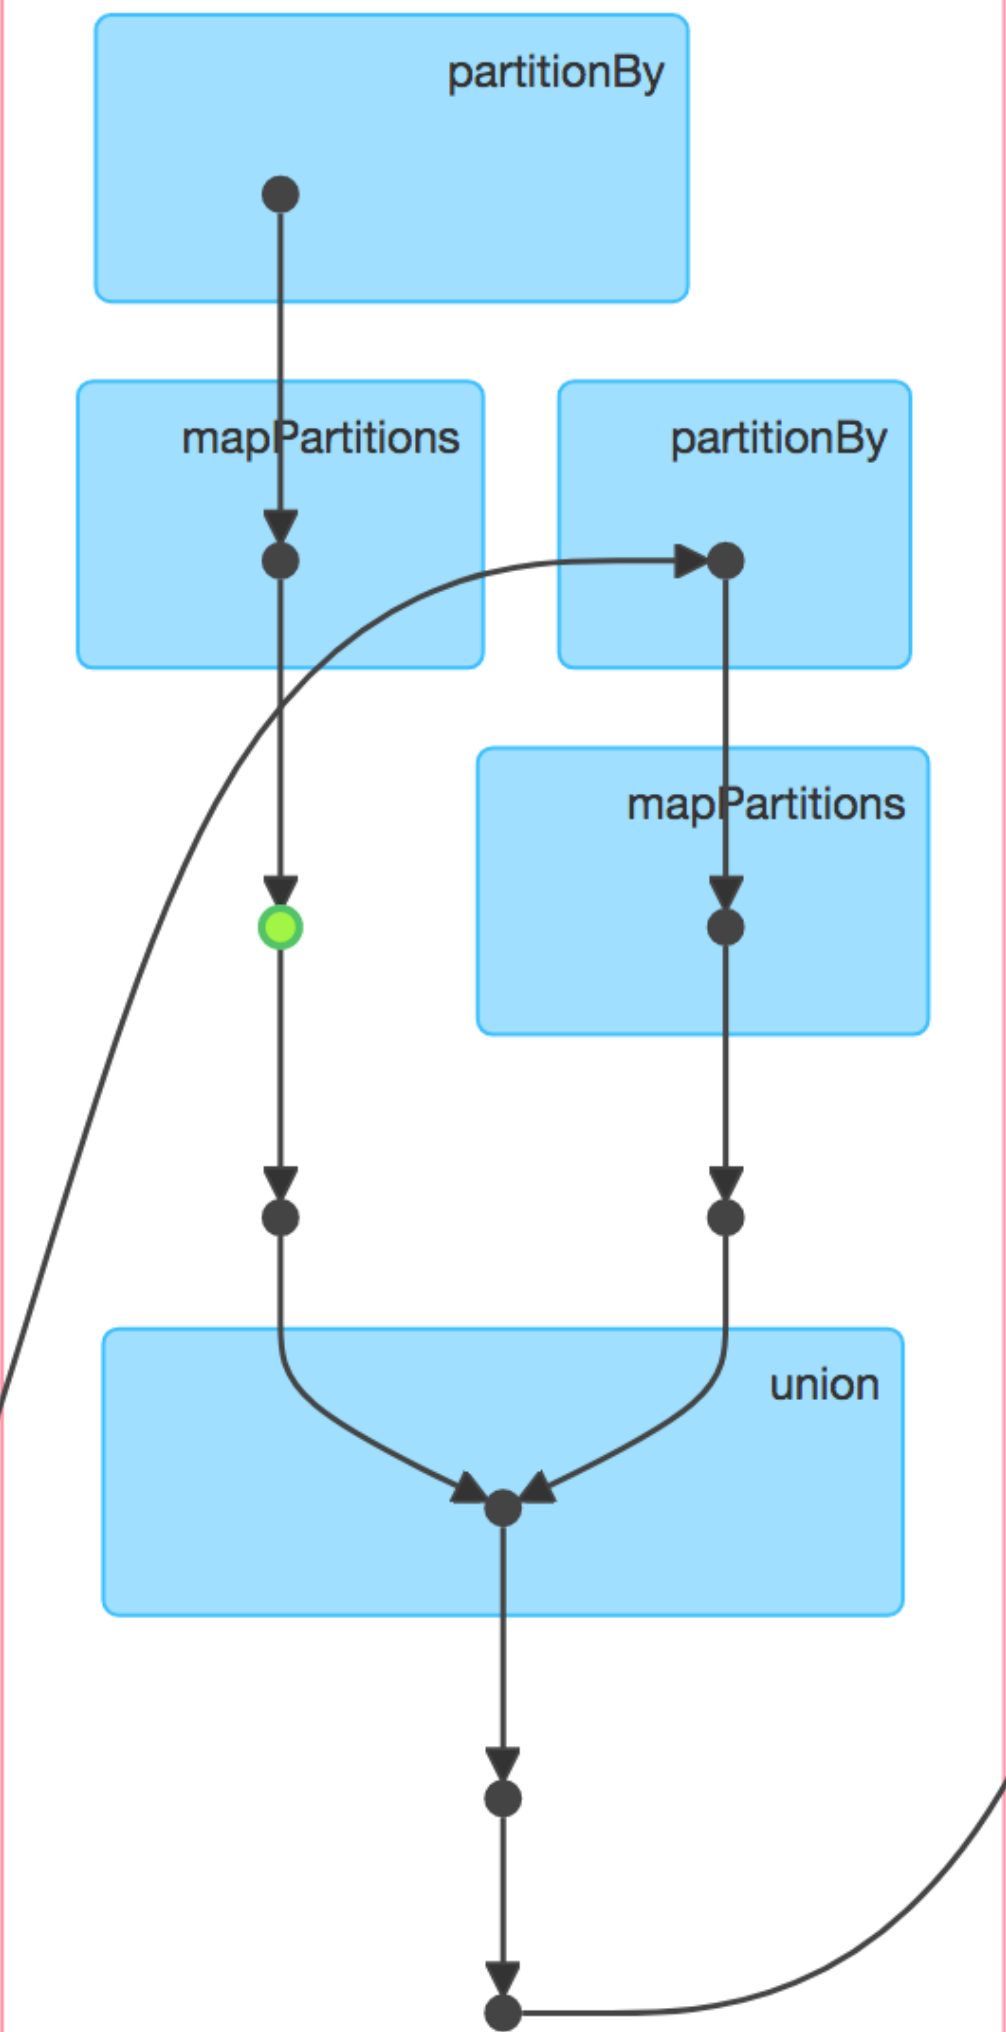
Completed Jobs (1)

Job Id ▾	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
0	collect at <magic-timeit>:257	2017/07/25 05:16:55	12 s	33/33	66/66

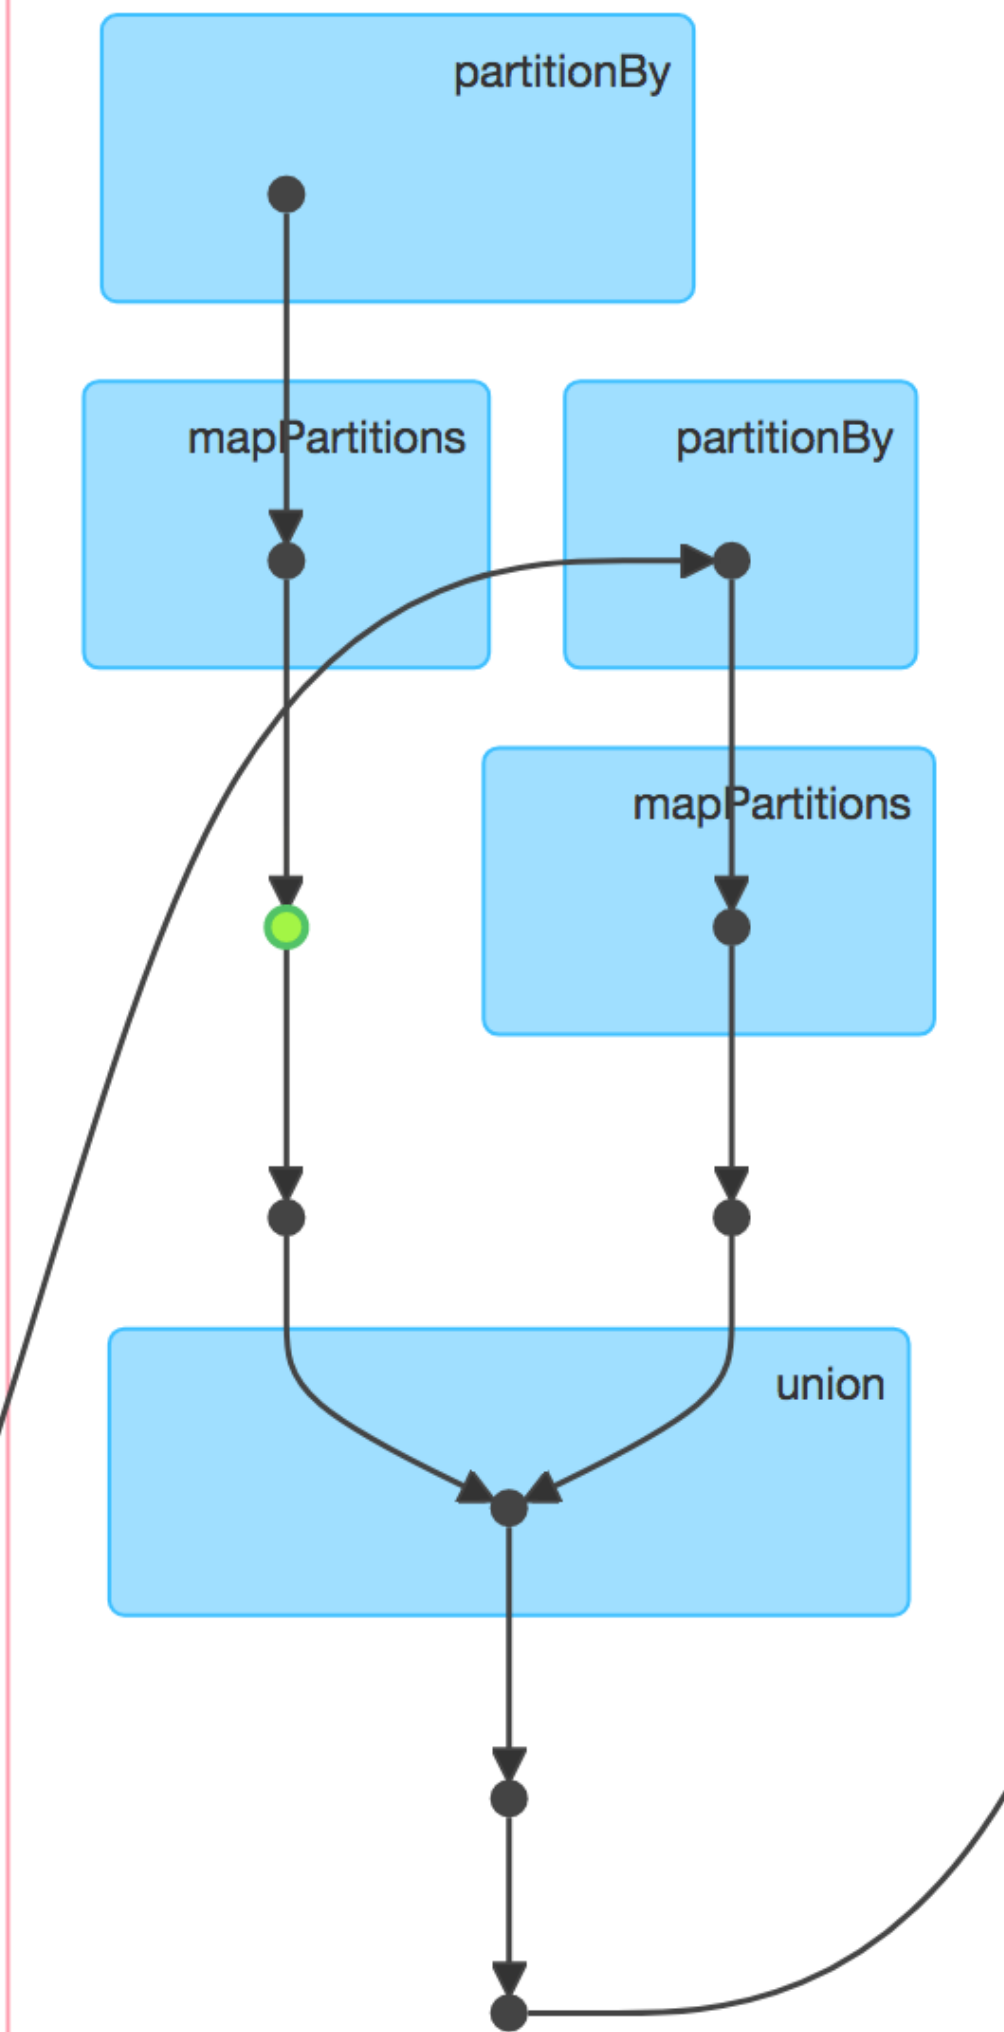
Stage 3



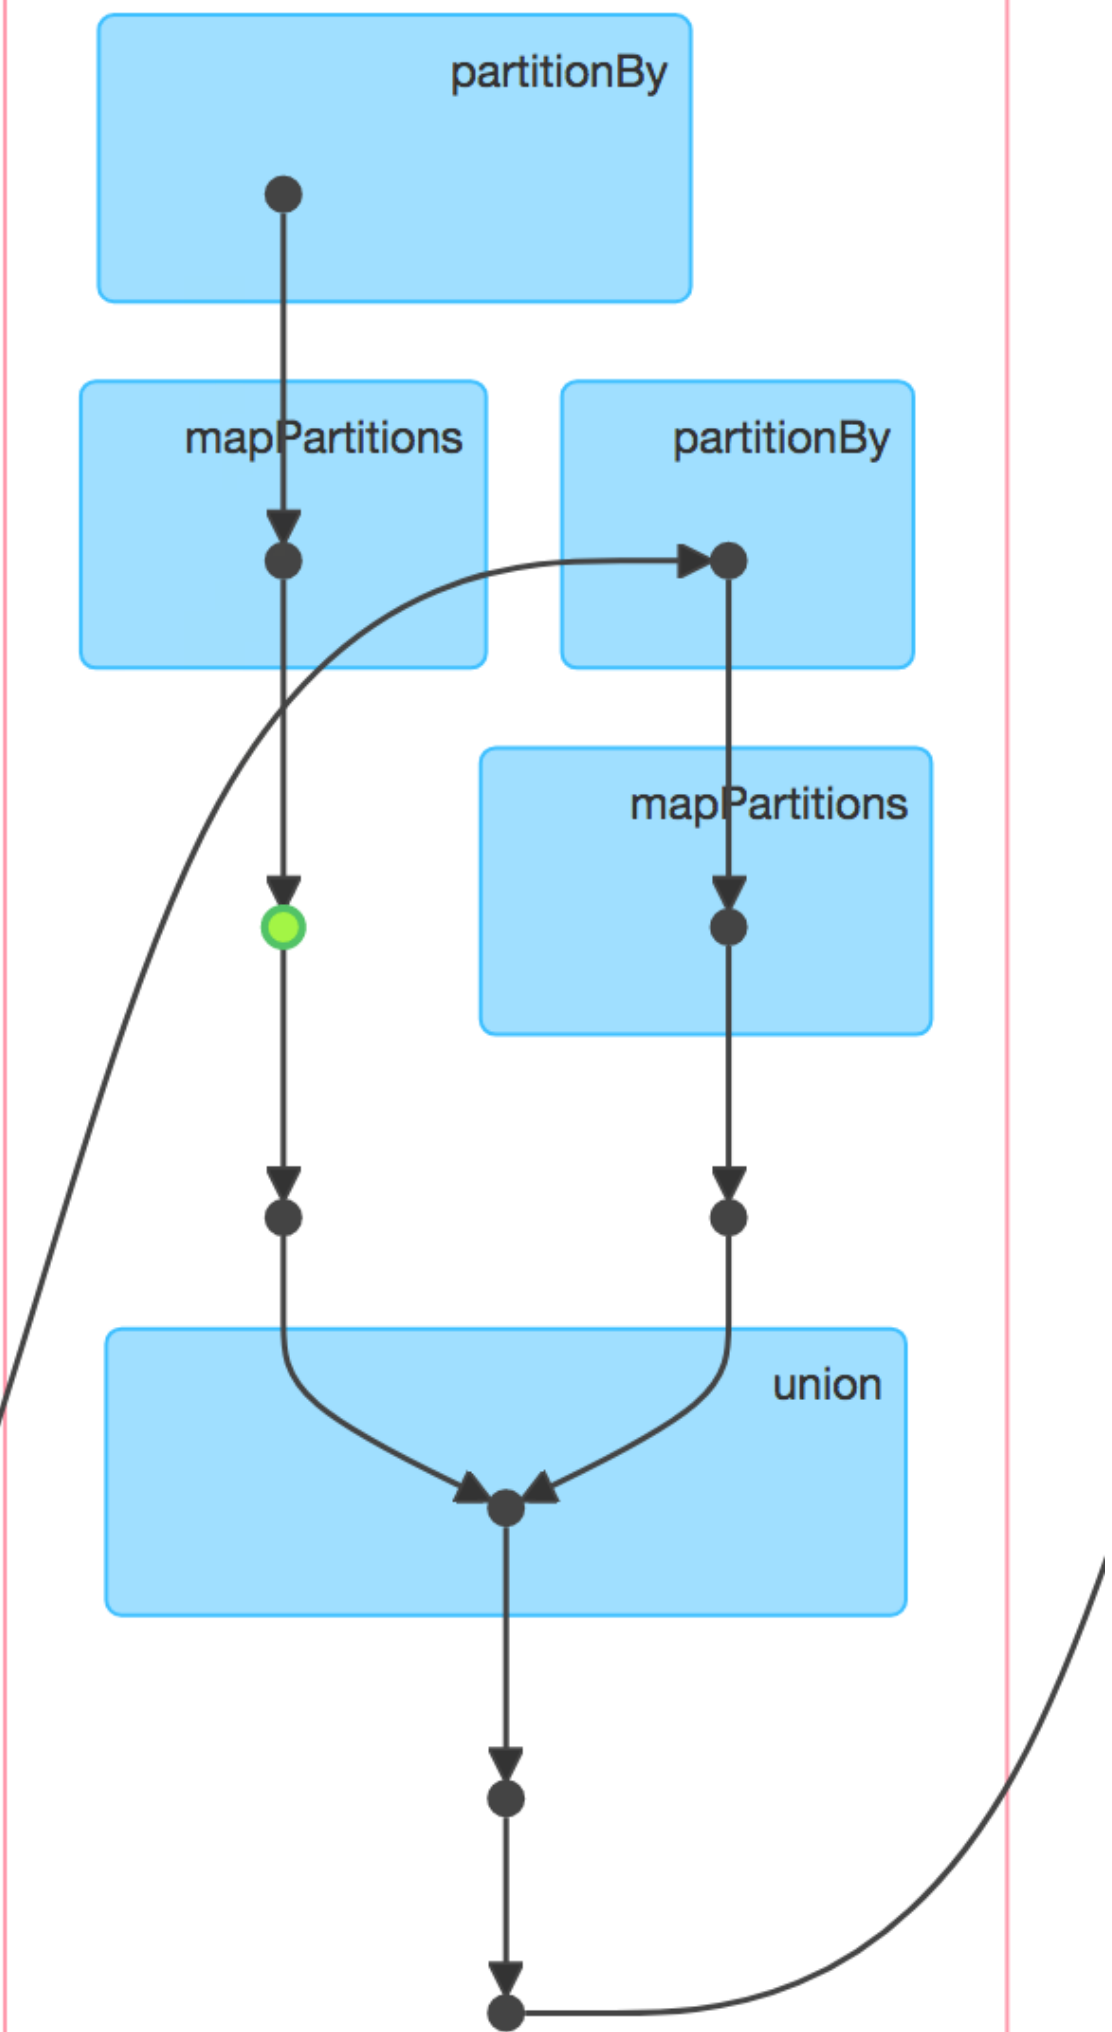
Stage 4



Stage 5



Stage 6



- 10 executors, 2 cores & 4GB per executor
- Berkeley-Stanford web graph
 - 685230 nodes
 - 7600595 edges
 - 106 MB uncompressed
- 5 iterations

Spark Jobs (?)

User: pklemenkov

Total Uptime: 5.6 min

Scheduling Mode: FIFO

Completed Jobs: 1

▶ [Event Timeline](#)

Completed Jobs (1)

Job Id ▼	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
0	collect at <magic-timeit>:257	2017/07/25 05:30:57	4.1 min	13/13	192/192

Spark Jobs (?)

User: pklemenkov

Total Uptime: 5.6 min

Scheduling Mode: FIFO

Completed Jobs: 1

▶ [Event Timeline](#)

Completed Jobs (1)

Job Id ▾	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
0	collect at <magic-timeit>:257	2017/07/25 05:30:57	4.1 min	13/13	192/192

~20% performance gain

Summary

- Reduce shuffles!
- Explicit (known) partitioner
- Preserve partitioner
- Reduce shuffle volume
 - `reduceByKey` vs `groupByKey`